

Chapter-5

December 22, 2018

0.1 Deep Learning in Computer Vision

Topics Covered:

1. Building Image Classifier on MNIST Dataset
2. Working of Conv in PyTorch
3. Why Pooling?
4. Why use View?
5. Transfer Learning
6. Visualizing Outputs from Intermediate Layers
7. Visualizing Weights of Intermediate Layers

```
In [1]: # Necessary Imports
import torch, torchvision
import time
import numpy as np
import matplotlib.pyplot as plt
from torch.autograd import Variable
import torchvision.transforms as transforms
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
```

0.1.1 Challenges of using Linear Layers/FCN Layers?

- Loss of Spatial Information.
 - Flattening an image into 1D array loses all the spatial information about the image.
- High complexity (number of weights for all layers in FCN)

0.1.2 Building an Image Classifier on MNIST Dataset

Steps:

1. Getting the Data
2. Pre-process (flattening, resizing -- if required, split the dataset - train_test_split)
3. Build model (CNN)
4. Train and Validate the model.

5. Test on Unseen Data (Test Dataset)

Step-1: Getting the Data

```
help(torchvision.datasets)
```

Output:

Help on package torchvision.datasets in torchvision:

NAME

torchvision.datasets

PACKAGE CONTENTS

cifar
coco
fakedata
folder
lsun
mnist
omniglot
phototour
semeion
stl10
svhn
utils

It's clear that we have mnist dataset available in torchvision datasets utility class. Let's go ahead and import mnist dataset.

```
In [5]: import torchvision.datasets.mnist as mnist
```

```
In [9]: # required transformations
```

```
transformation = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.1307, 0.3081), (0.3081, 0.3081))])
```

```
In [11]: train_dataset = mnist.MNIST('data/', train=True, transform=transformation, download=True)
test_dataset = mnist.MNIST('data/', train=False, transform=transformation, download=True)
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz>

Processing...

Done!

```
In [14]: train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=32, shuffle=True)
```

```
In [36]: # utility function to convert a tensor into an image
def show_image(tensor, cmap=None):
    image = tensor.numpy()[0]
    mean = 0.1307
    std_dev = 0.3081
    image = ((mean * image) + std_dev)
    plt.imshow(image, cmap) # show gray-scaled version of the image
    plt.show()

In [29]: next(iter(train_loader))[0] # returns batch of 32 images, randomly

Out[29]: tensor([[[[-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  ...,
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242]]],

  [[[-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  ...,
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242]]],

  [[[-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  ...,
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242]]],

  ...,

  [[[-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  ...,
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
  [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242]]]]]
```

```

[[-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
 [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
 [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
 ...,
 [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
 [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
 [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242]]],

[[-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
 [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
 [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
 ...,
 [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
 [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242],
 [-0.4242, -0.4242, -0.4242, ..., -0.4242, -0.4242, -0.4242]]]]))


```

```


In [39]: print("Showing image in color mode")
         show_image(next(iter(train_loader))[0][0])
         print("Showing image in grayscale mode")
         show_image(next(iter(train_loader))[0][0], cmap='gray')

```

Showing image in color mode



Showing image in grayscale mode



0.2 Building Model

Generally, a CNN is composed of following layers:

1. Conv2d
2. MaxPooling2d
3. ReLU
4. View
5. Linear Layer (FC Layer)
6. Dropout

```
In [42]: class Net(nn.Module):
          def __init__(self):
              super().__init__()
              self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
              self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
              self.conv2d_dropout = nn.Dropout2d()
              self.fc1 = nn.Linear(320, 50)
              self.fc2 = nn.Linear(50, 10)
          def forward(self, x):
              x = F.relu(F.max_pool2d(self.conv1(x), 2))
              x = F.max_pool2d(self.conv2d_dropout(self.conv2(x)), 2)
              x = x.view(-1, 320) # why?
```

```

x = F.relu(self.fc1(x))
x = F.dropout(x, training=self.training)
x = self.fc2(x)
return F.log_softmax(x)

```

0.3 How does Convolution works in PyTorch?

Take an example of a 1D array. Let's try to apply a conv filter on the array.

```
conv = nn.Conv1d(1, 1, 3, bias=False)
```

Let's look at the docs of `nn.Conv1d`:

```

class Conv1d(_ConvNd)
|   Applies a 1D convolution over an input signal composed of several input
|   planes.
|
|   In the simplest case, the output value of the layer with input size
|   (N, C_{in}, L) and output (N, C_{out}, L_{out}) can be
|   precisely described as:

```

$$\text{out}(N_i, C_{out_j}) = \text{bias}(C_{out_j}) + \sum_{k=0}^{C_{in}-1} \text{weight}(C_{out_j}, k) \star \text{input}(N_i, k)$$

```

|   Attributes:
|       weight (Tensor): the learnable weights of the module of shape
|           (out_channels, in_channels, kernel_size)
|       bias (Tensor): the learnable bias of the module of shape
|           (out_channels)
|
|   Examples::
|
|       >>> m = nn.Conv1d(16, 33, 3, stride=2)
|       >>> input = torch.randn(20, 16, 50)
|       >>> output = m(input)

```

```

In [76]: net = nn.Conv1d(16, 33, 3, stride=2)
print("Net.weight: ", net.weight[0][0])
print("Net bias: ", net.bias[0][0])
input_ = torch.randn(20, 16, 50)
print("Shape of input tensor: {}".format(input_.shape))
output = net(input_)
print("Input", input_[0][0])
print("output", output[0][0])

```

```

Net.weight: tensor([ 0.0898, -0.0290, -0.0184], grad_fn=<SelectBackward>)
Net bias: tensor(-0.0234, grad_fn=<AliasBackward>)
Shape of input tensor: torch.Size([20, 16, 50])

```

```

Input tensor([ 0.4412,  1.1871,  1.2538,  0.9038, -0.4481,  1.3991, -0.9111,  0.2678,
               1.1055,  1.9549, -1.1615, -1.0567,  0.5989,  0.1253,  0.7138,  1.2872,
               1.2061, -0.4276, -1.0316, -2.1270,  0.6886,  0.5826, -1.3476,  0.6016,
              -0.3757, -1.1891,  0.6825, -0.0293, -1.0270,  0.8466,  1.6183, -1.1577,
              -1.6786,  1.4142, -0.6939, -0.8775, -1.0439,  0.1629,  1.2870, -1.2488,
              -0.8653,  0.8719,  0.2489, -2.0293, -2.2702,  0.8735,  0.9727,  0.5769,
              -0.4099, -0.9260])
output tensor([-0.4045,  0.4693, -0.6951, -0.4687, -0.0146, -0.4754,  0.5161, -0.1440,
              -0.2339, -0.2184,  0.3562, -0.5512, -0.7846, -0.8485, -0.1358,  0.3537,
              -0.2544,  0.1960, -0.3285,  0.1912, -0.1416,  0.0757,  0.1543,  1.5208],
              grad_fn=<SelectBackward>)

```

/home/kushashwa/.local/lib/python3.6/site-packages/ipykernel_launcher.py:3: UserWarning: invalid operation:
This is separate from the ipykernel package so we can avoid doing imports until

```
In [64]: net.weight[0][0]
```

```
Out[64]: tensor([-0.0673,  0.1119, -0.0590], grad_fn=<SelectBackward>)
```

```
In [79]: net.weight[0][0]
```

```
Out[79]: tensor([ 0.0898, -0.0290, -0.0184], grad_fn=<SelectBackward>)
```

```
In [91]: net.bias[0]
```

```
Out[91]: tensor(-0.0234, grad_fn=<SelectBackward>)
```

```
In [88]: print(input_[0][0][:3] * net.weight[0][0])
         print("input_, net.weight, net.bias: {}/{}{}".format(input_[0][0][:3], net.weight[0][0], net.bias[0]))
```

```
tensor([ 0.0396, -0.0344, -0.0231], grad_fn=<ThMulBackward>)
input_, net.weight, net.bias: tensor([0.4412, 1.1871, 1.2538])/tensor([ 0.0898, -0.0290, -0.0184])/tensor(-0.0234)
```

/home/kushashwa/.local/lib/python3.6/site-packages/ipykernel_launcher.py:2: UserWarning: invalid operation:
This is separate from the ipykernel package so we can avoid doing imports until

```
In [104]: np.dot(input_[0][0][:3].detach().numpy(), net.weight[0][0].detach().numpy()) + net.bias[0].detach().numpy()
```

```
Out[104]: -0.041310795
```

```
In [98]: type(net.weight)
```

```
Out[98]: torch.nn.parameter.Parameter
```

0.4 Why Pooling?

Used just after a convolution layer, to reduce the data size to process. Also helps in reducing the size of feature maps. Also, forces algorithm to not focus on small changes in position. (**how?**)

0.5 Why View?

Generally, we use `torch.Tensor.view()` (<https://pytorch.org/docs/stable/tensors.html#torch.Tensor.view>) function at the end of a network, because for FC (Fully Connected) or Linear layers, we need to flatten the data to 1D.

While flattening, it's important to make sure that two different images don't mix-up. That's why we input the first argument to `torch.Tensor.view()` as `-1`.

Let's look at the docs of this function.

```
In [166]: print(help(torch.Tensor.view))
```

Help on method_descriptor:

```
view(...)
```

```
view(*args) -> Tensor
```

Returns a new tensor with the same data as the `:attr:`self`` tensor but of a different size.

The returned tensor shares the same data and must have the same number of elements, but may have a different size. For a tensor to be viewed, the new view size must be compatible with its original size and stride, i.e., each new view dimension must either be a subspace of an original dimension, or only span across original dimensions `:math:`d, d+1, \dots, d+k`` that satisfy the following contiguity-like condition that `:math:`\forall i = 0, \dots, k-1``,

```
.. math::
```

```
stride[i] = stride[i+1] \times size[i+1]
```

Otherwise, `:func:`contiguous`` needs to be called before the tensor can be viewed.

Args:

args (torch.Size or int...): the desired size

Example::

```
>>> x = torch.randn(4, 4)
>>> x.size()
torch.Size([4, 4])
>>> y = x.view(16)
>>> y.size()
torch.Size([16])
>>> z = x.view(-1, 8) # the size -1 is inferred from other dimensions
>>> z.size()
torch.Size([2, 8])
```


None

```
In [164]: torch.Tensor.view?
```

```
In [158]: x = np.array([[[[8], [9]], [[2], [0.3]]], [[8], [9]], [[2], [0.3]]])  
          x = x.reshape((2, 1, 2, 2)) # reshape to (2, 1, 2, 2) - like a batch of 2 images of .
```

```
In [159]: x.view() # view the array x
```

```
Out[159]: array([[[[8. , 9. ],  
                  [2. , 0.3]]],  
                  
                [[8. , 9. ],  
                 [2. , 0.3]]])
```

```
In [160]: x_tensor = torch.from_numpy(x) # conversion of numpy array to a pytorch tensor  
          # reference: https://pytorch.org/tutorials/beginner/blitz/tensor\_tutorial.html#conve
```

```
In [161]: print(x_tensor.shape) # verifying
```

```
torch.Size([2, 1, 2, 2])
```

```
In [167]: x_tensor.view(-1, 4)  
          # -1 means: don't touch the first dimension (only if the second dimension satisfies  
          # if doesn't satisfy, then -1 is inferred from other dimensions
```

```
Out[167]: tensor([[8.0000, 9.0000, 2.0000, 0.3000],  
                  [8.0000, 9.0000, 2.0000, 0.3000]], dtype=torch.float64)
```

```
In [171]: x_tensor.view(-1, 2) # example: this should return (4, 2) tensor
```

```
Out[171]: tensor([[8.0000, 9.0000],  
                  [2.0000, 0.3000],  
                  [8.0000, 9.0000],  
                  [2.0000, 0.3000]], dtype=torch.float64)
```

```
In [174]: x_tensor.view(-1, 1) # example: this should return (8, 1) tensor
```

```
Out[174]: tensor([[8.0000],  
                  [9.0000],  
                  [2.0000],  
                  [0.3000],  
                  [8.0000],  
                  [9.0000],  
                  [2.0000],  
                  [0.3000]], dtype=torch.float64)
```

0.6 Training the Model

```
In [197]: def fit(epoch, optimizer, model, data_loader, volatile=False, phase='training'):
    if(phase == 'training'):
        model.train()
    if(phase == 'evaluation'):
        model.evaluate()
        volatile=True #why?

    running_loss = 0.0
    running_correct = 0

    for batch_idx, (data, target) in enumerate(data_loader):
        #         print(data)
        #         if data.is_cuda():
        #             data, target = data.cuda(), target.cuda()
        data, target = Variable(data, volatile), Variable(target)
        if phase == 'training':
            optimizer.zero_grad()
            output = model(data)
            loss = F.nll_loss(output, target)
            running_loss += F.nll_loss(output, target, size_average=False).data[0]
            preds = output.data.max(dim=1, keepdim=True)[1]
            running_correct += preds.eq(target.data.view_as(preds)).cpu().sum()
        if phase == 'training':
            loss.backward()
            optimizer.step()

    loss = running_loss/len(data_loader.dataset)
    accuracy = 100. * running_correct/len(data_loader.dataset)
    print("Loss: {}, accuracy: {}".format(loss, accuracy))
    return loss, accuracy

In [198]: model = Net()
          optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.5)

In [199]: optimizer

Out[199]: SGD (
  Parameter Group 0
    dampening: 0
    lr: 0.01
    momentum: 0.5
    nesterov: False
    weight_decay: 0
)

In [200]: train_losses, train_accuracy = [], []
          val_losses, val_accuracy = [], []
```

```
In [204]: for epoch in range(1, 20):
            epoch_loss, epoch_accuracy = fit(epoch, optimizer, model, train_loader, volatile_device_ids)
            val_epoch_loss, val_epoch_accuracy = fit(epoch, optimizer, model, test_loader, volatile_device_ids)
            train_losses.append(epoch_loss)
            train_accuracy.append(epoch_accuracy)
            val_losses.append(val_epoch_loss)
            val_accuracy.append(val_epoch_accuracy)
```


```
/home/kushashwa/.local/lib/python3.6/site-packages/ipykernel_launcher.py:16: UserWarning: Impl
app.launch_new_instance()
/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:52: UserWarning: size_average and
warnings.warn(warning.format(ret))
/home/kushashwa/.local/lib/python3.6/site-packages/ipykernel_launcher.py:20: UserWarning: inva
```

```
Loss: 0.1206219345331192, accuracy: 96
Loss: 0.12402940541505814, accuracy: 96
Loss: 0.11611782014369965, accuracy: 96
Loss: 0.11210104078054428, accuracy: 96
Loss: 0.11929728090763092, accuracy: 96
Loss: 0.1225995272397995, accuracy: 96
Loss: 0.11394266784191132, accuracy: 96
Loss: 0.11298033595085144, accuracy: 96
Loss: 0.11729313433170319, accuracy: 96
Loss: 0.11354488879442215, accuracy: 96
Loss: 0.115715891122818, accuracy: 96
Loss: 0.11816882342100143, accuracy: 96
Loss: 0.11266960203647614, accuracy: 96
Loss: 0.11292938143014908, accuracy: 96
Loss: 0.11479900777339935, accuracy: 96
Loss: 0.12624315917491913, accuracy: 96
Loss: 0.11431062966585159, accuracy: 96
Loss: 0.1180247887969017, accuracy: 96
Loss: 0.1105051189661026, accuracy: 96
Loss: 0.12324003875255585, accuracy: 96
Loss: 0.11057820171117783, accuracy: 96
Loss: 0.10865706950426102, accuracy: 97
Loss: 0.10824849456548691, accuracy: 96
Loss: 0.11275696009397507, accuracy: 96
Loss: 0.10923698544502258, accuracy: 96
Loss: 0.11321546882390976, accuracy: 96
Loss: 0.10587245225906372, accuracy: 96
Loss: 0.11644107848405838, accuracy: 96
Loss: 0.10572720319032669, accuracy: 96
Loss: 0.10592672973871231, accuracy: 96
Loss: 0.10624512284994125, accuracy: 96
Loss: 0.12218598276376724, accuracy: 96
Loss: 0.1034180149435997, accuracy: 96
```

```
Loss: 0.114189513027668, accuracy: 96
Loss: 0.10513007640838623, accuracy: 96
Loss: 0.11694205552339554, accuracy: 96
Loss: 0.10253959894180298, accuracy: 96
Loss: 0.10903823375701904, accuracy: 96
```

```
In [205]: plt.plot(range(1,len(train_accuracy)+1),train_accuracy,'bo',label = 'train accuracy')
          plt.plot(range(1,len(val_accuracy)+1),val_accuracy,'r',label = 'val accuracy')
          plt.legend()
```

```
Out[205]: <matplotlib.legend.Legend at 0x7fa4d4396940>
```



0.7 Transfer Learning

- Reusing pre-trained algorithm on a dataset similar to what we are using. Training from scratch not required.
- Freeze most of the layers and fine tune params for some of the layers only.

```
In [3]: from torchvision import models
```

```
In [208]: vgg = models.vgg16(pretrained=True) # load the pretrained model VGG16
```

```
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /home/kushashwa/.torch
100%|| 553433881/553433881 [00:18<00:00, 30166213.80it/s]
```

```
In [210]: # let's look at the VGG network
         print(vgg)
```

```
VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace)
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): ReLU(inplace)
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (18): ReLU(inplace)
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (20): ReLU(inplace)
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (22): ReLU(inplace)
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (25): ReLU(inplace)
    (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (27): ReLU(inplace)
    (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (29): ReLU(inplace)
    (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (classifier): Sequential(
    (0): Linear(in_features=25088, out_features=4096, bias=True)
    (1): ReLU(inplace)
    (2): Dropout(p=0.5)
    (3): Linear(in_features=4096, out_features=4096, bias=True)
    (4): ReLU(inplace)
    (5): Dropout(p=0.5)
    (6): Linear(in_features=4096, out_features=1000, bias=True)
  )
)
```

In the above summary of VGG16 Model, there are 2 Sequential models:

1. Features: - Has the layers that we are going to freeze **2. Classifier:** - Has the layers that we are going to learn.

```
In [211]: # freeze the weights for features layers
          for param in vgg.features.parameters():
              param.requires_grad = False
              # prevent optimizer from updating the weights

In [217]: # let's fine tune the classifier layers
          # for dogs and cats?
          vgg.classifier[6].out_features = 2

In [218]: optimizer = optim.SGD(vgg.classifier.parameters(), lr=0.001, momentum=0.5)
          # only pass classifier parameters for optimization
```

Quick Tips to improve accuracy

```
In [229]: # changing dropout from 0.2 to 0.5
          for layer in vgg.classifier.children():
              if(type(layer) == nn.modules.dropout.Dropout):
                  layer.p = 0.5

In [230]: # performing data augmentation to add more data: flipping, mirroring, rotating with
          train_transform = transforms.Compose([transforms.Resize((224, 224)),
                                                transforms.RandomHorizontalFlip(),
                                                transforms.RandomRotation(0.2),
                                                transforms.ToTensor(),
                                                transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])])
```

0.8 Visualizing Outputs from Intermediate Layers

We can visualize outputs from Intermediate Layers, using PyTorch utility function: `register_forward_hook`

```
In [4]: vgg = models.vgg16(pretrained=True)

In [5]: class LayerActivations():
          features=None

          def __init__(self,model,layer_num):
              self.hook = model[layer_num].register_forward_hook(self.hook_fn)

          def hook_fn(self,module,input,output):
              self.features = output.cpu().data.numpy()

          def remove(self):
              self.hook.remove()

In [6]: conv_out = LayerActivations(vgg.features, 0)
```

Let's create a sample image

```
In [52]: img = []  
        for i in range(64):  
            img.append(np.random.randn(16, 16, 3))
```

```
In [53]: img = np.array(img)  
        img.shape
```

```
Out[53]: (64, 16, 16, 3)
```


This is the image created.

```
In [54]: img = img.reshape((64, 16, 16, 3))
```

```
In [55]: plt.imshow(img[0])
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255]

```
Out[55]: <matplotlib.image.AxesImage at 0x7f5cf7f02ba8>
```



```
In [56]: img = img.reshape((64, 3, 16, 16))
```

```
In [57]: img = torch.from_numpy(img).float()
```

```
In [58]: print(img.shape)
```

```
torch.Size([64, 3, 16, 16])
```

This wasn't working on my PC, lack of free RAM you know? ;)

Thanks to Google Collabs :D Implemented it here: (on a random input image)

<https://colab.research.google.com/drive/12UH6rswmxIuvIvvHT-wjB0LNVj-Txw9>

0.9 Visualizing weights of Intermediate Layers

```
In [65]: cnn_weights = vgg.state_dict()['features.0.weight'].cpu()
```

```
In [68]: # directly taken from the book
fig = plt.figure(figsize=(30,30))
fig.subplots_adjust(left=0,right=1,bottom=0,top=0.8,hspace=0,wspace=0.2)
for i in range(30):
    ax = fig.add_subplot(12,6,i+1,xticks=[],yticks=[])
    plt.imshow(cnn_weights[i])
```

[illegible]

